Discrete Geometry for Risk and AI

Dr. Paul Larsen

March 15, 2022

Why discrete geometry?

- Recent history: Dissatisfaction with deep learning, only "curve fitting", alternatives via *causal graphical models* [**?**]
- Less recent history: graphical models among first non-rules based AI approaches [**?**]
- Geometrical formulations of statistical objects, e.g. graphical models and probability polytopes

Directed graphical model: university admission gender bias

Simpson paradox preview

	Men		Women	
	Applicants	**Admitted**	**Applicants**	**Admitted**
Total	8442	**44%**	4321	35%

Directed graphical model: university admission gender bias

Simpson paradox preview

Department	Men		Women	
	Applicants	Admitted	Applicants	Admitted
A	*825*	62%	108	**82%**
B	*560*	63%	25	**68%**
C	325	**37%**	*593*	34%
D	417	33%	375	**35%**
E	191	**28%**	*393*	24%
F	373	6%	341	**7%**

Sources: [**?**] [**?**]

Directed graphical model: hit rate for insurance quotes

- product type: financial, liability, property
- days: number of days to generate quote
- rating: measure of premium paid expected claims
- hit: 0 if quote refused, 1 if accepted

Undirected graphical model: credit default risk [?]

- Nodes take values 0 (healthy) or 1 (default)
- Industry nodes connect to other industry nodes
- Individual firm nodes connect only to corresponding industry node

Graph definitions

Definition
A *graph* is a pair of sets (V, E), where V is called the set of *vertices* (or *nodes*) and E is called the set of *edges*, such that the set of edges corresponds injectively to pairs of vertices.

Notes
- Typically 'pairs of vertices' does not include self-pairs, but this can be relaxed, leading to graphics with with loops.
- The injectivity requirement can also be relaxed, leading to *multigraphs*.

Graphical models

Definition
(Informal) A graphical model is a graph whose nodes represent variables and whose edges represent direct statistical dependencies between the variables.

Why graphical models?

- For probability distributions admitting a graphical model representation, then graph properties (*d-separation*) imply conditional independence relations.
- Conditional independence relations reduce the number of parameters required to specify a probability distribution.
- Graphical models come in two flavors depending on their edges: directed (aka *Bayesian Networks*) and undirected (aka *random Markov fields*.

Directed acyclic graphs

Definition
A graph $G = (V, E)$ is a *directed acyclic graph* (denoted also *DAG*) if all edges have an associated direction, and no edge path consistent with the directions forms a cycle.

If there is a directed path from X_i to X_j, then X_i is called a *parent* of X_j, and $Pa(X_j) \subseteq V$ is the set of all parents of X_j.

Definition
If $X = (X_1, \ldots, X_m)$ admits a DAG G, then X_G is a *DAG model* if the distribution of X decomposes according to G, i.e.

$$P(X) = \prod_{i \in \{1, \ldots, m\}} P(X_i | Pa(X_i))$$

Example: Karma and weight-lifting

Take K to be your Karma, H to be the hours you spend in the gym lifting weight each day, and then W be the weight you can bench press on a given day. For simplicity, all random variables are binary.

karma	hours	weight
1	0	1
1	1	1
0	1	0
1	0	1
1	0	1

Decomposition example: Karma and weight-lifting

Suppose $X = (K, H, W)$ admits the graph

Then $P(K, H, W) = P(K) \, P(H|K) P(W|H)$.

Definition

A DAG of the form above is called a *chain*.

Decomposition example: Karma and weight-lifting

Suppose $X = (K, H, W)$ admits the graph

Then $P(K, H, W) = P(K)\, P(H)\, P(W|K, H)$.

Definition

A DAG of the form above is called a *collider* at W.

Conditional independence

Recall that two random variables X, Y are *independent* if
$P(X = x, Y = y) = P(X = x)P(Y = y)$.

Definition

Let $X = (X_1, \ldots, X_m)$ be a probability distribution, and let A, B, C be pair-wise disjoint subsets of $1, \ldots, m$, and define $X_A = (X_i)_{i \in A}$. Then X_A, X_B are *conditionally depenedent given X_C* if and only if

$$
P(X_A = x_A, X_B = x_B | X_C = x_c) \\
= P(X_A = x_a | X_C = x_c)P(X_B = x_B | X_C = x_c)
$$

for all x_A, x_B, x_C.

For X_A, X_B conditionally independent given X_C, we write $(X_A \perp\!\!\!\perp X_B | X_C)$. See e.g. [?] for a precise formulation.

Conditional independence and d-separation teaser

First example of discrete geometry helping statistics: conditional independence in a DAG model (X, G) can be detected in properties of G[1]. More precisely,

Theorem
If (X, G) is a DAG model, then d-separation implies conditional independence.

See e.g. [**?**], chapter 2.

---

[1]The required graph properties are combinatorial, but can also be understood geometrically, see e.g. [**?**].

More definitions before d-separation

Figure: Chain

Figure: Collider at W, Fork at K

d-separation in DAGs

Definition

An undirected path p in a DAG G is *blocked* by a set of nodes C if and only if

1. p contains a chain of nodes $X \rightarrow Y \rightarrow Z$, or a fork $X \leftarrow Y \rightarrow Z$ such that $Y \in C$, or

2. p contains a collider $X \rightarrow Y \leftarrow Z$ such that $Y \notin C$ and descendant of Y is in C.

Definition

If C blocks every path between two nodes X and Y, then X and Y are called *d-separated conditional on* C, and we write

$$(X \perp\!\!\!\perp Y | C)_G$$

.

By the d-separation teaser theorem, $(X \perp\!\!\!\perp Y | C)_G$ implies conditional independence.

d-separation example: hit rate for insurance

All paths from $\mathrm{product\_type}$ to hit are blocked by $\{\mathrm{days}, \mathrm{rating}\}$, hence $(\mathrm{product\_type} \perp\!\!\!\perp \mathrm{hit}|\mathrm{days}, \mathrm{rating})_G$.

Probability polytopes

Goal: Use geometric interpretation of multivariate discrete random variables to generate intersting fake data with few(er) parameters.

Example: The family of all $X \sim$ *Bernoulli* can be represented as

$$\Delta_1 = \{(p_0, p_1) : p_i \geq 0, \sum p_i = 1\} \subseteq R^2$$

Example: Consider the collider graph for Karma-influenced weight-lifting (K, H, W). Then all possible conditional probability tables for $(W|K, H)$ can be parametrized as

$$\{(p_{w|k,h}) : p_{w|k,h} \geq 0, \sum_w p_{w|k,h} = 1 \text{ for } (k, h) \in \{0, 1\}^2\} \subseteq R^8$$

In general, the space of multivariate discrete random variable distributions is a *polytope*, see e.g. [**?**], Ch. 1.

H- and V-representations of polytopes

Definition
An *H-polyhedron* is an intersection of closed halfspaces, i.e. a set $P \subseteq R^d$ presented in the form

$$P = P(A, z) = \{x \in R^d : Ax \leq z\} \text{ for some } A \in R^{md}, z \in R^m.$$

If P is bounded (i.e. compact), then it is called a *polytope*.

Definition
(Informal) A *V-polytope* is the convex hull of a finite set of vertices $conv(V) \in R^d$.
See [?] for a precise definition.

Example: The V-representation for all *Bernoulli* distributions is

The main theorem of polytopes

Theorem
A subset $P \subset R^d$ is the convex hull of a finite point set (a V-polytope)

$$P = conv(V) \; for \; some \; V \in R^{dn}$$

if and only if it is a bounded intersection of halfspaces (an H-polytope)

$$P = P(A, z) \; for \; some \; A \in R^{md}, z \in R^m$$

See [?] for a proof.

Applying the main theorem to conditional probability tables

For the Karma weight-lifting example, all conditional probability tables for $(W|K, H)$ that satisfy $E(W|K = 0) = 0$ (bad Karma, no weight) and $E(W|H = 0) = 0.2$ can be written as an $H - polytope$ as above with additional constraints

$$\sum_{w,h} w \, p_{w|0,h} = 0$$

$$\sum_{w,k} w \, p_{w|k,0} = 0.2$$

By converting this H-representation to a V-representation, we can generate random conditional probability tables subject to expectation constrains.

For an example, see the implementation of ProbabilityPolytope of https://munichpavel.github.io/fake-data-for-learning/.

References I